



Are You Ready?

Lets Begin the Unlimited Learning

Operations on Resources

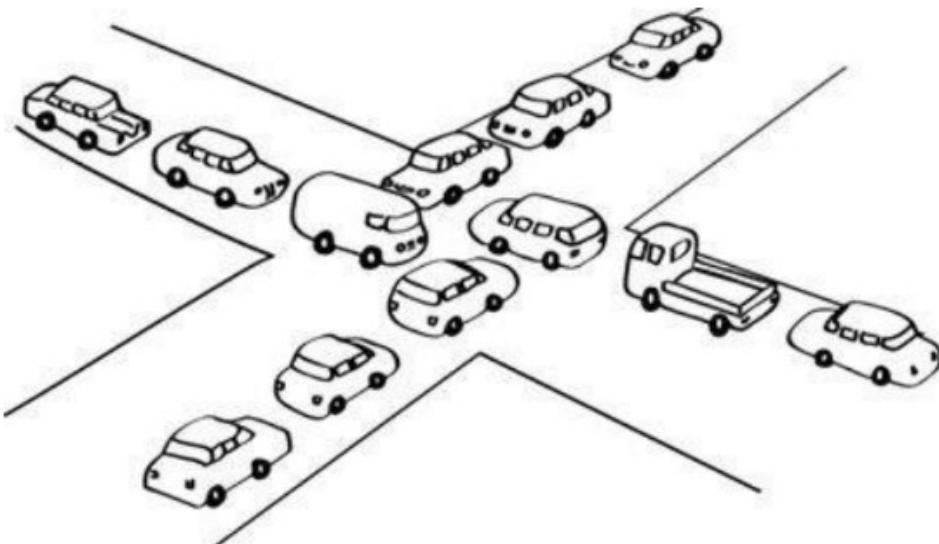
3 operations on resources:

1. Request ✓
2. Use ✓
3. Release ✓



Deadlock

If two or more processes are waiting for such an event which is never going to occur



Necessary Conditions for Deadlock

Deadlock can occur only when all following conditions are satisfied:

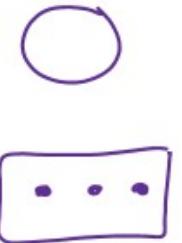
1. Mutual Exclusion
2. Hold & Wait
3. No-preemption
4. Circular Wait



Resource Allocation Graph

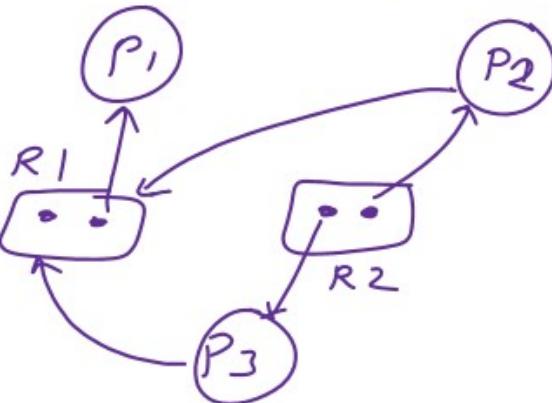
Nodes

- 1. Process
- 2. Resource



Edges

- Allocation from instance to process
- Request —||— process to resource



unacademy

Recovery From Deadlock

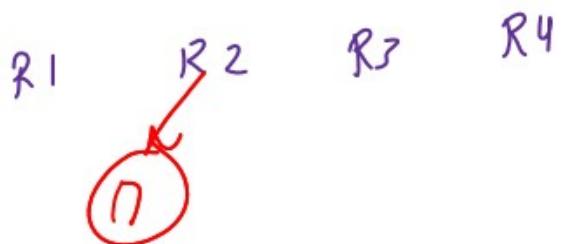
- 
1. Make Sure that deadlock never occur
 - Prevent the system from deadlock or avoid deadlock
 2. Allow deadlock, detect and recover
 3. Pretend that there is no any deadlock



Deadlock Prevention

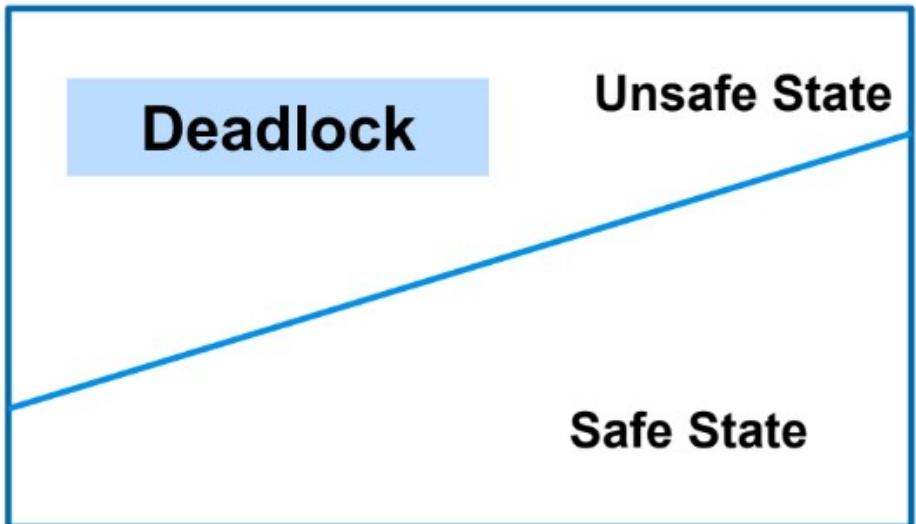
Prevent any of four necessary conditions to occur

1. Mutual Exclusion
2. Hold & Wait
3. No Preemption
4. Circular Wait



Deadlock Avoidance

In deadlock avoidance, the OS tries to keep system in safe state



Banker's Algorithm

The banker's algorithm is a resource allocation and deadlock avoidance algorithm that tests for safety

→ safety
→ resource request



unacademy

Banker's Algorithm

Process	Allocation			Max			Available			Need
	A	B	C	A	B	C	A	B	C	
P ₀	0	1	0	7	5	3	3	3	2	7 4 3
P ₁	2	0	0	3	2	2	1	5	2	2 2
P ₂	3	0	2	9	0	2	1	7	4	3
P ₃	2	1	1	2	2	2	0	0	0	1 1
P ₄	0	0	2	4	3	3	0	1	1	4 3 1

Safe
seq. = P₁, P₃, P₂, P₀, P₄

Resource Request Algorithm

1. $\text{Req}_i \leq \text{Need}_i$ ✓

2. $\text{Req}_i \leq \text{Available}$

3. $\text{Alloc}^n_i = \text{Alloc}^r_i + \text{Req}_i$

$$\text{Available} = \text{Available} - \text{Req}_i$$

$$\text{Need}_i = \text{Need}_i - \text{Req}_i$$

4. Safety Algo → safe \rightarrow granted
 unsafe \rightarrow rejected

Question

What will happen if process P1 requests one additional instance of resource type A and two instances of resource type C?

$\text{Req}_1 <1, 0, 2>$



Banker's Algorithm

Process	Allocation			Max			Available		
	A	B	C	A	B	C	A	B	C
P ₀	0	1	0	7	5	3	3	3	2
P ₁	3	2	0	2	0	2	2	3	0
P ₂	3	0	2	9	0	2			
P ₃	2	1	1	2	2	2			
P ₄	0	0	2	4	3	3			

$$\begin{array}{r} +2 \\ \hline 0 & 2 & 0 \end{array}$$

Safe → Granted

Recovery From Deadlock

1. Make Sure that deadlock never occur
 - Prevent the system from deadlock or avoid deadlock
2. Allow deadlock, detect and recover
3. Pretend that there is no any deadlock



Deadlock Detection

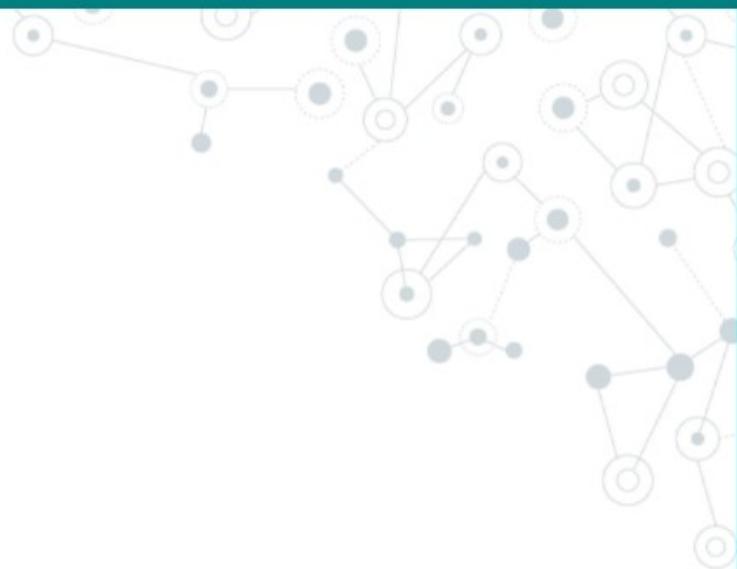
1. When all resources have single instance
2. When resources have multiple instances \Rightarrow Banker's algo



Deadlock Detection

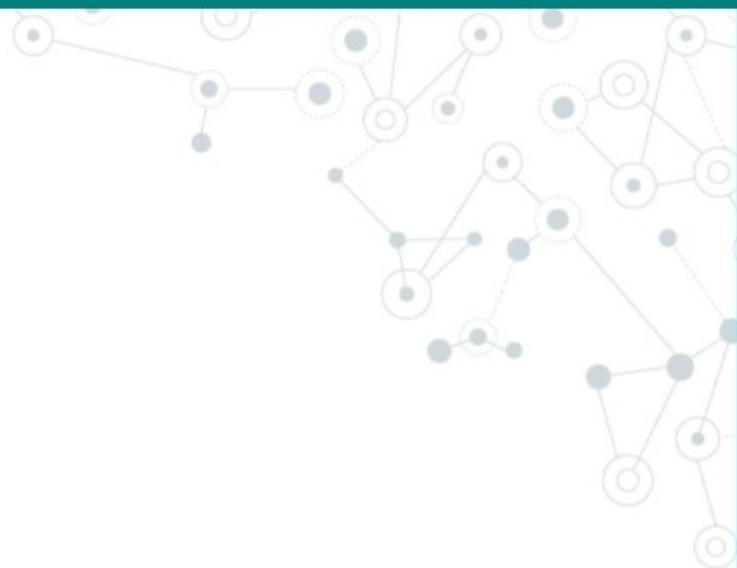
When all resources have single instance:

Deadlock detection is done using wait-for-graph

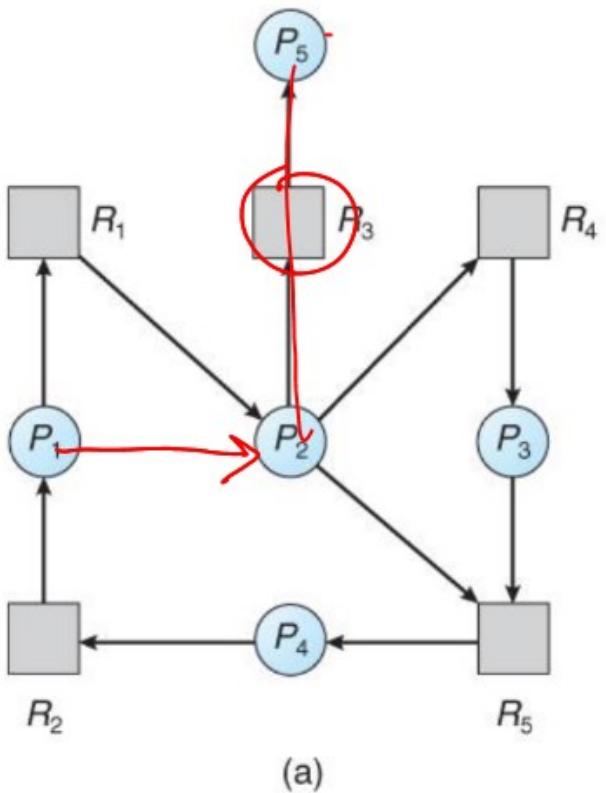


Wait For Graph

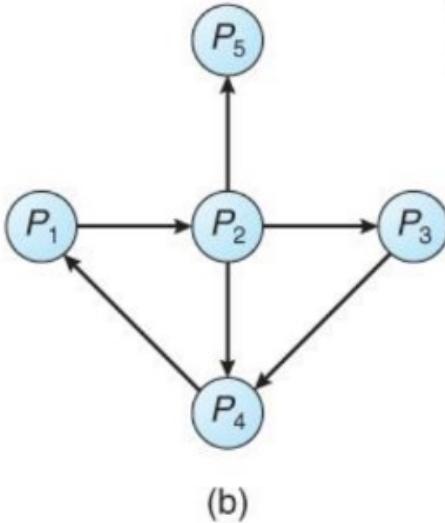
It is created from resource allocation graph



Wait For Graph

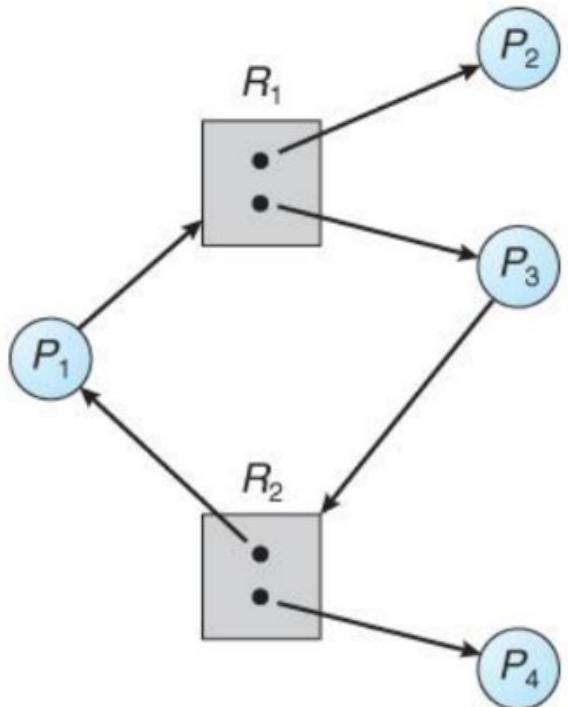


(a)



(b)

Wait For Graph: Example

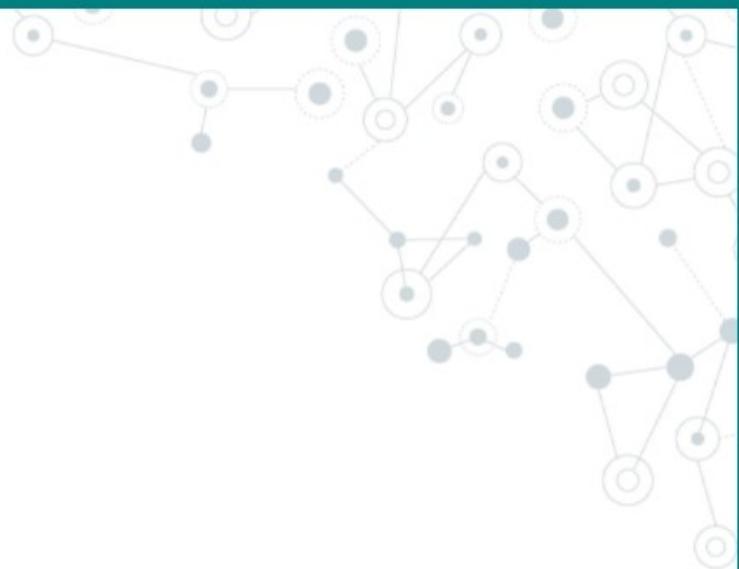


		Allocat ^h		Reg.	
		R1	R2	R1	R2
P1		0	1	1	0
	P2	1	0	0	0
P3	1	0	0	1	
P4	0	1	0	0	

Deadlock Detection

When resources have multiple instance:

Deadlock detection is done using a specific algorithm



Deadlock Detection Algorithm

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	0 1 0
P_2	3 0 3	0 0 0	3 1 3
P_3	2 1 1	1 0 0	5 1 3
P_4	0 0 2	0 0 2	7 3 4
			7 3 6



Question

Consider a system with 3 processes A, B and C. All 3 processes require 4 resources each to execute. The minimum number of resources the system should have such that deadlock can never occur? 10

n processes

k resources

$$n(k-1) + 1$$



Recovery From Deadlock

There are three basic approaches to recovery from deadlock:

1. Inform the system operator and allow him/her to take manual intervention
2. Terminate one or more processes involved in the deadlock
3. Preempt resources.

Process Termination

1. Terminate all processes involved in the deadlock
2. Terminate processes one by one until the deadlock is broken



Process Termination

Many factors that can go into deciding which processes to terminate next:

1. Process priorities.
2. How long the process has been running, and how close it is to finishing.
3. How many and what type of resources is the process holding
4. How many more resources does the process need to complete
5. How many processes will need to be terminated
6. Whether the process is interactive or batch

Resource Preemption

Important issues to be addressed when preempting resources to relieve deadlock:

1. Selecting a victim
2. Rollback
3. Starvation



Memory Management

- ◎ Module of OS
- ◎ Functions:
 1. Memory allocation
 2. Memory deallocation
 3. Memory protection
- ◎ Goals:
 1. Maximum Utilization of space
 2. Ability to run larger programs with limited space



Memory Management Techniques

Contiguous

Fixed partition

Variable partition

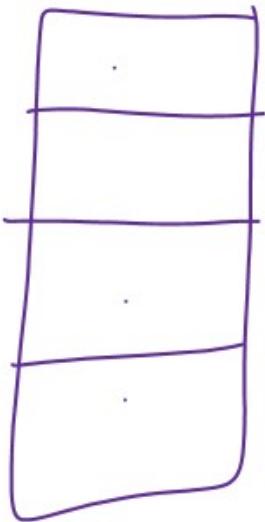
Non-contiguous

Paging

Segmentation



Fixed Partition Contiguous MMT

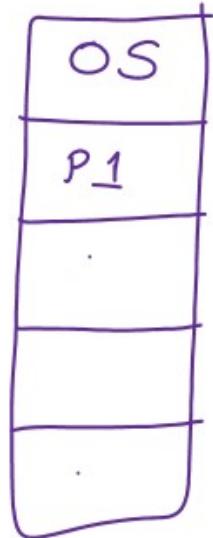


Partition Allocation Policy

- First fit
- Best fit
- Worst fit



Variable Partition Contiguous MMT



External Fragmentation }
Compaction



Non-Contiguous MMT

- ◎ Process is scattered in memory, not allocated at one area
- ◎ Two techniques:
 - Paging: Scattered in same size of memory areas
 - Segmentation: Scattered in variable size of memory areas



Paging

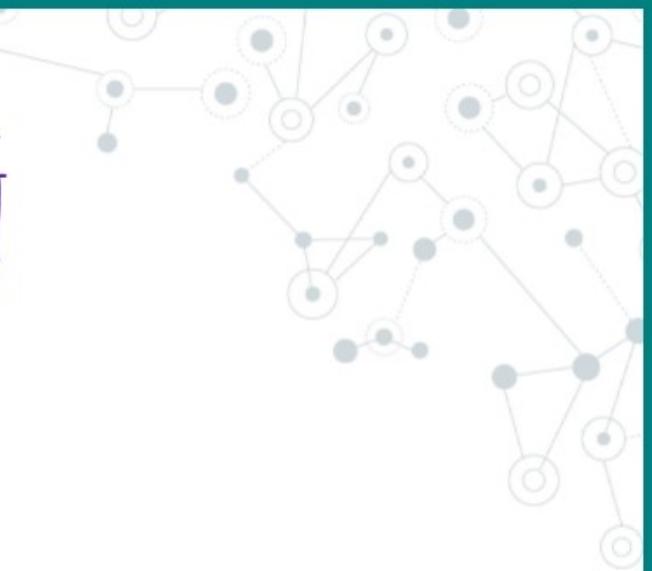
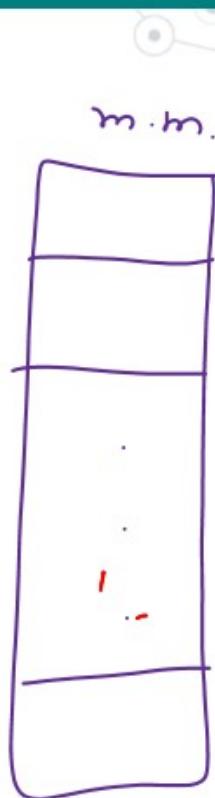
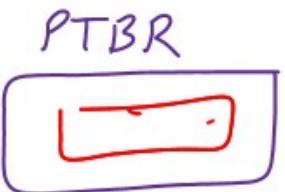
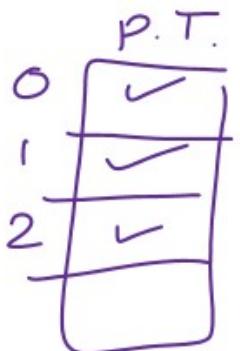
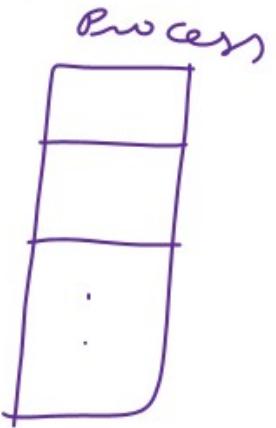
- Process is divided in equal size of pages
- Physical memory is divided in same equal size of frames
- Processor will have a view of process and its pages
- Pages are scattered in frames
- Page table is used to map a process page to a physical frame

1 P.T. entry = frame no. + protection bits
(translatⁿ)

No. of entries in PT = No. of pages in process

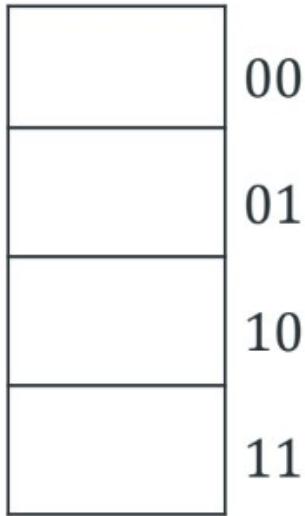


Paging: Example

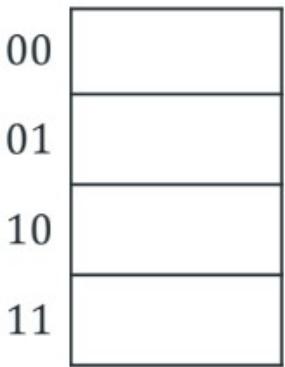


Paging

Process



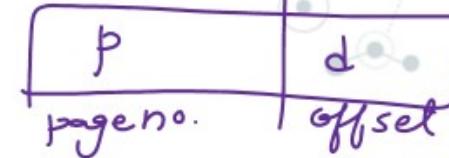
Page Table



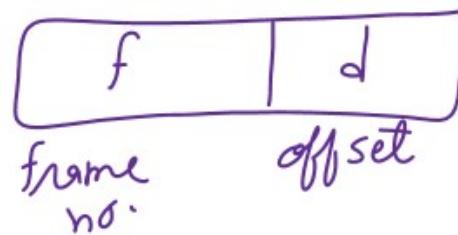
Physical Memory

0000	
0001	
0010	
0011	
0100	
0101	
0110	
0111	
1000	
1001	
1010	
1011	
1100	
1101	
1110	
1111	

L.A.



P.A.



Paging

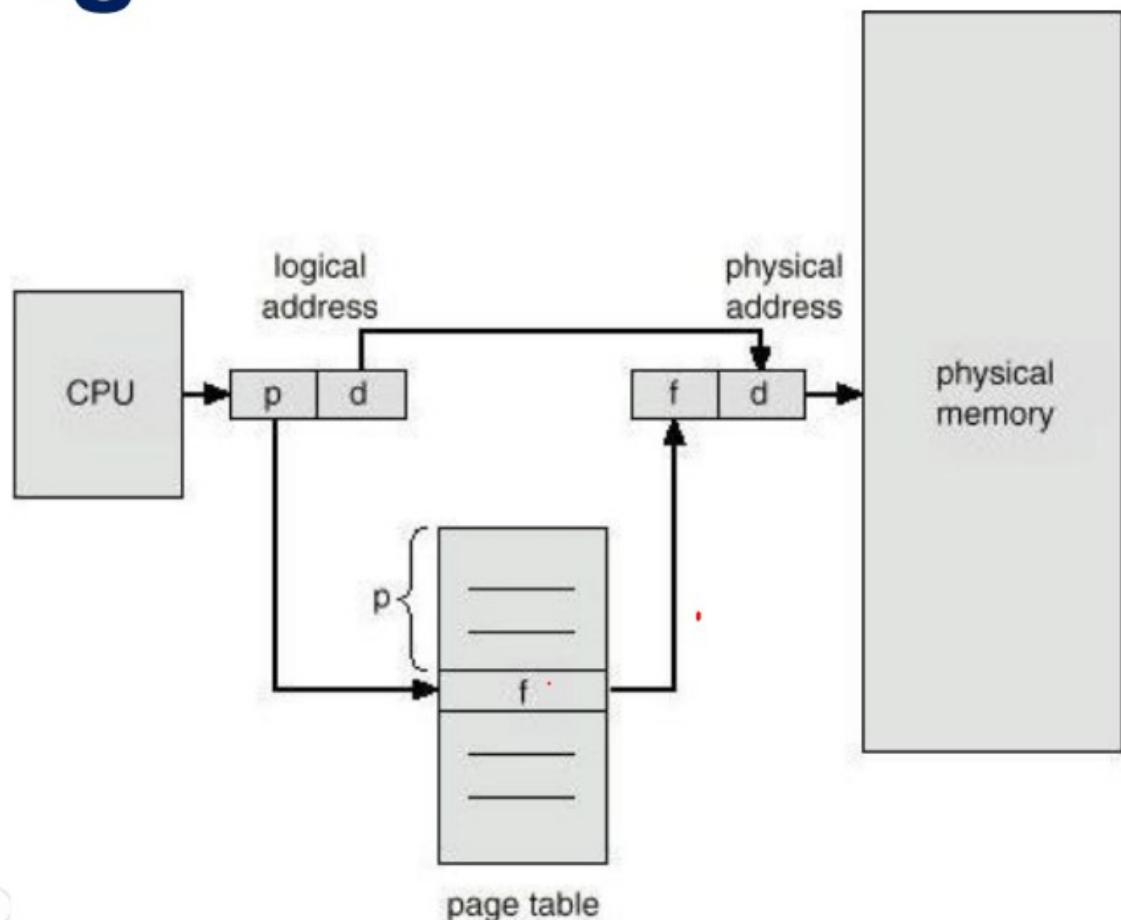
$$\begin{aligned} P.T. \text{ size} &= \text{no. of entry} * \text{1 entry size} \\ &= \text{no. of pages} * (\text{f} + \text{extra bits}) \end{aligned}$$

$$\text{no. of pages} = \frac{\text{L.A.S.}}{\text{Page size}}$$

$$\text{no. of frames} = \frac{\text{P.A.S.}}{\text{page size}}$$



Paging



Question 1

Consider a paged memory system with logical address of 28-bits and physical address of 33-bits. The page size is 2KB. Further consider that one page table entry size is 4bytes.

1. Bits in page offset $\Rightarrow 11$ bits
 2. Number of pages in process 2^{17}
 3. Bits for page number 17
 4. Number of frames in physical memory 2^{22}
 5. Bits for frame number 22
 6. Page table size $\hookrightarrow 2^{17} * 4B$
- extra bits in each P.T. entry = ? = 10 bits
- entry [f | extra]

Time Required in Paging

$$E.M.A.T. = 2 * t_{mm}$$

if P.T. is Registers

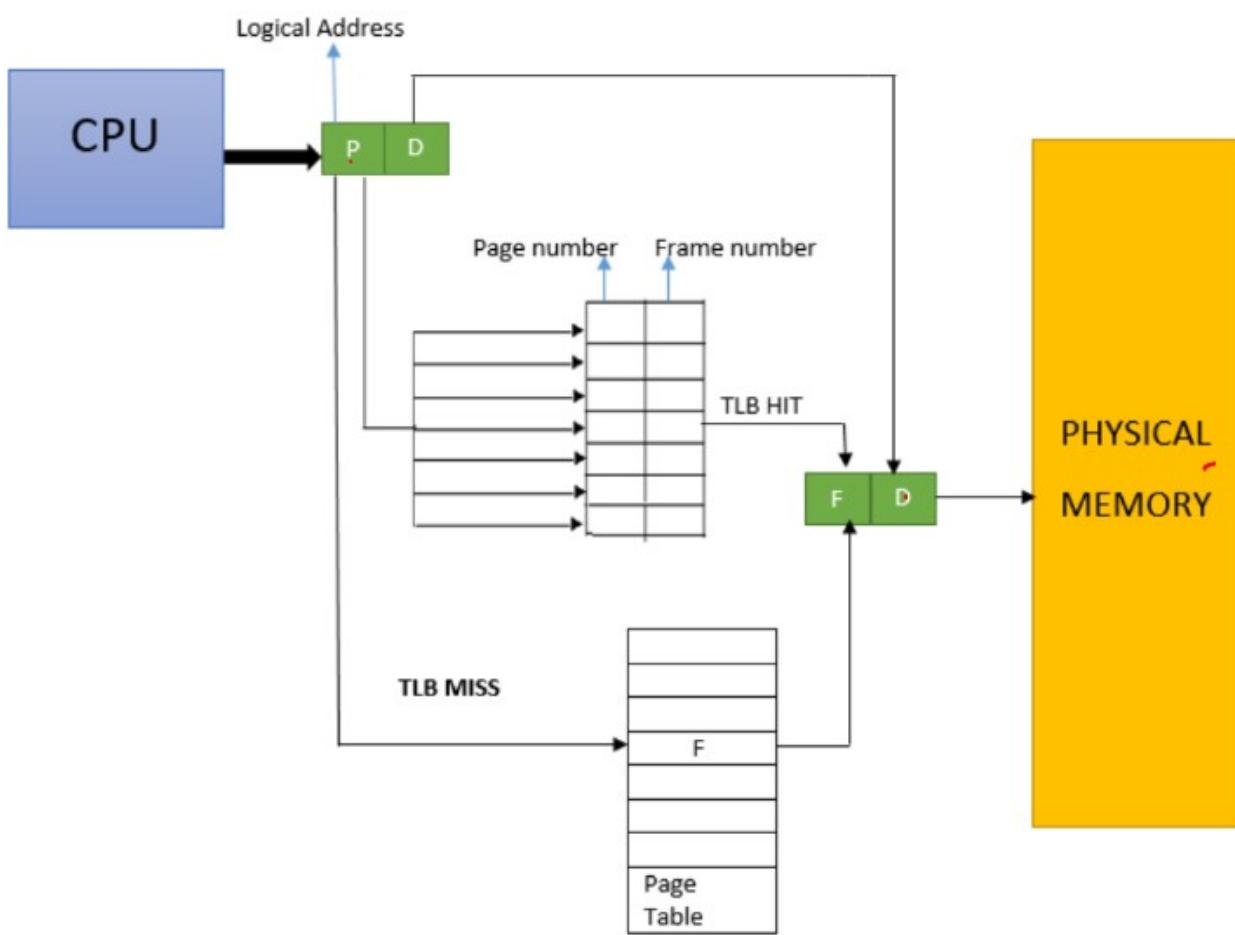
$$= t_{mm}$$

TLB

A translation lookaside buffer is a memory hardware that is used to reduce the time taken to access a user memory location



TLB



Effective Access Time With TLB

$$\text{E.m.A.T.} = t_{TLB} + t_{mm} + (1-H) t_{mm}$$

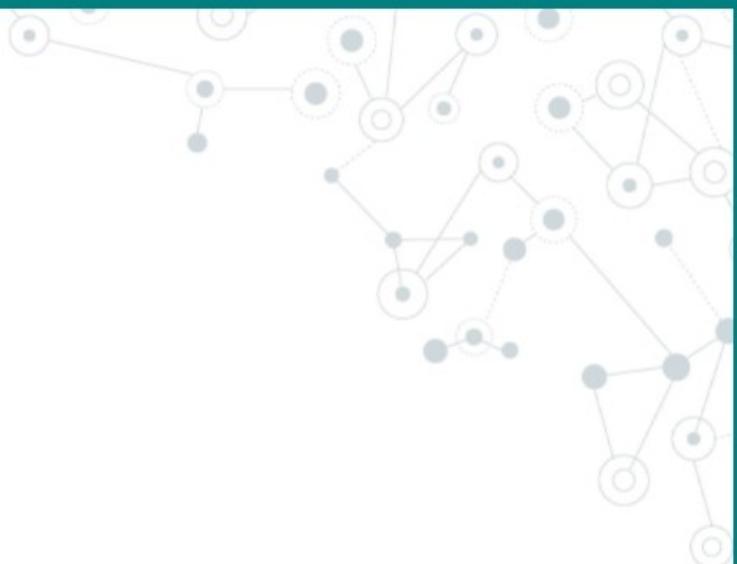


TLB

A translation lookaside buffer is a memory hardware that is used to reduce the time taken to access a user memory location



TLB Mappings

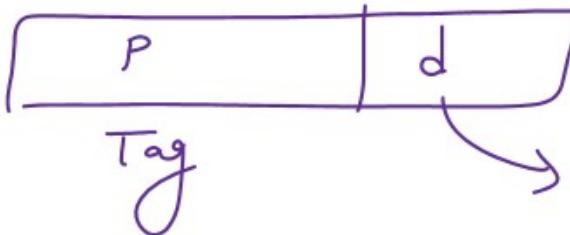


Full Associative Mappings

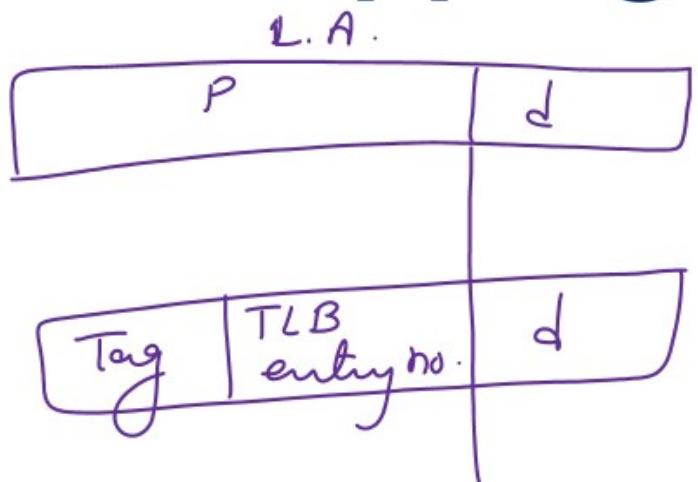
↓
Costly

P. No. (Tag)	F. no.
:	
:	
:	

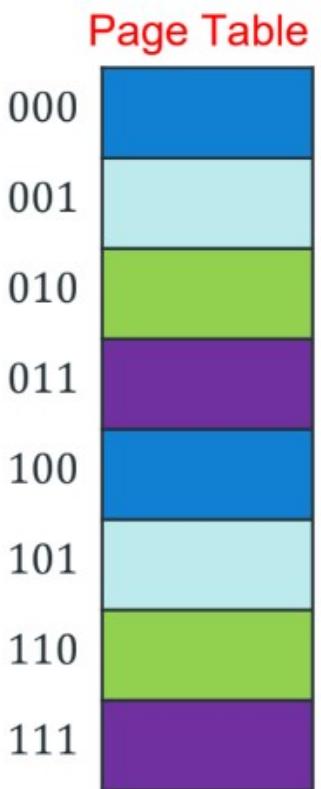
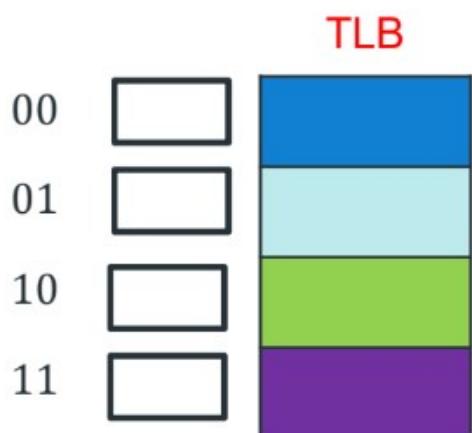
L.A.



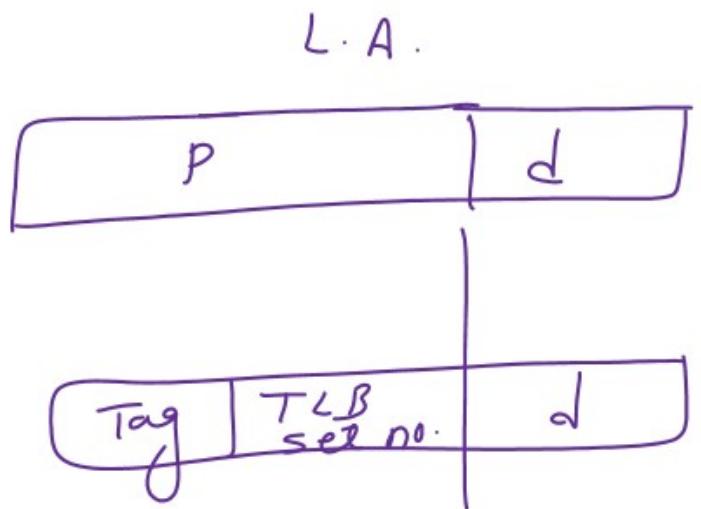
Direct Mappings



Problem With Direct Mappings



Set Associative Mappings



$$\frac{\text{no. of sets in TLB}}{\text{no. of entries in TLB}} = \text{associative}$$

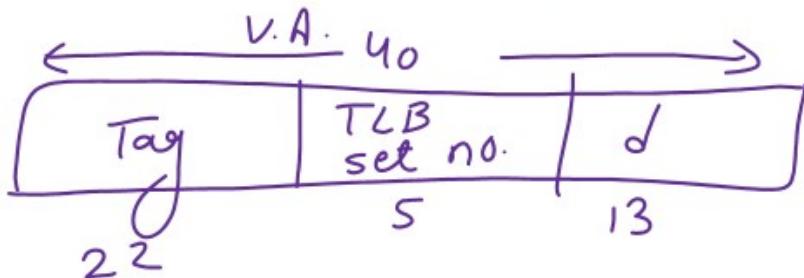


Question GATE-2015

A computer system implements a 40 bit virtual address, page size of 8 kilobytes, and a 128-entry translation look-aside buffer (TLB) organized into 32 sets each having four ways. Assume that the TLB tag does not store any process id. The minimum length of the TLB tag in bits is

_____?

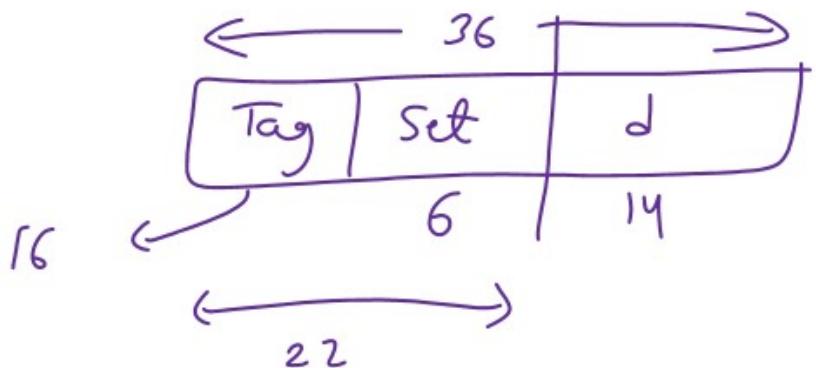
- (A) 20
- (B) 10
- (C) 11
- (D) 22



Question

Ans = 16

A Computer system implements a 36-bits virtual address, page size of 16Kbytes and a 256-entry TLB organized into 64 sets each having 4-ways. Assume that the TLB tag does not store any process-id. The minimum length of the TLB tag in bits is _____?

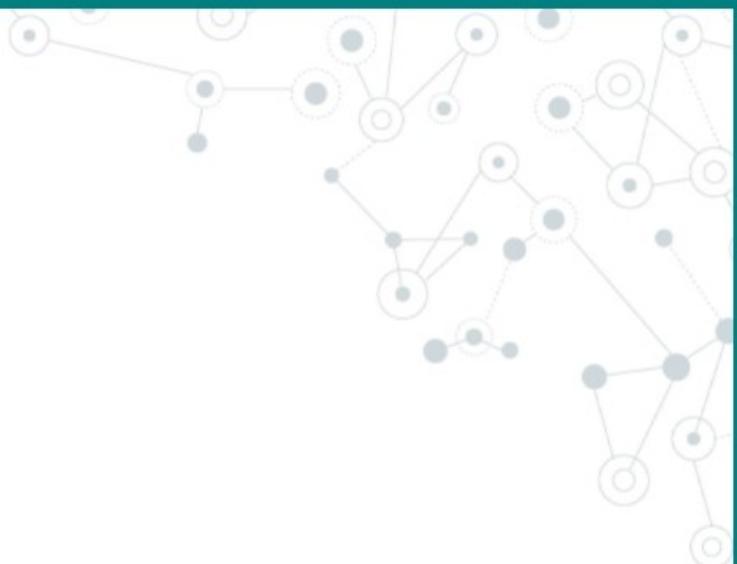


Optimal Page Size

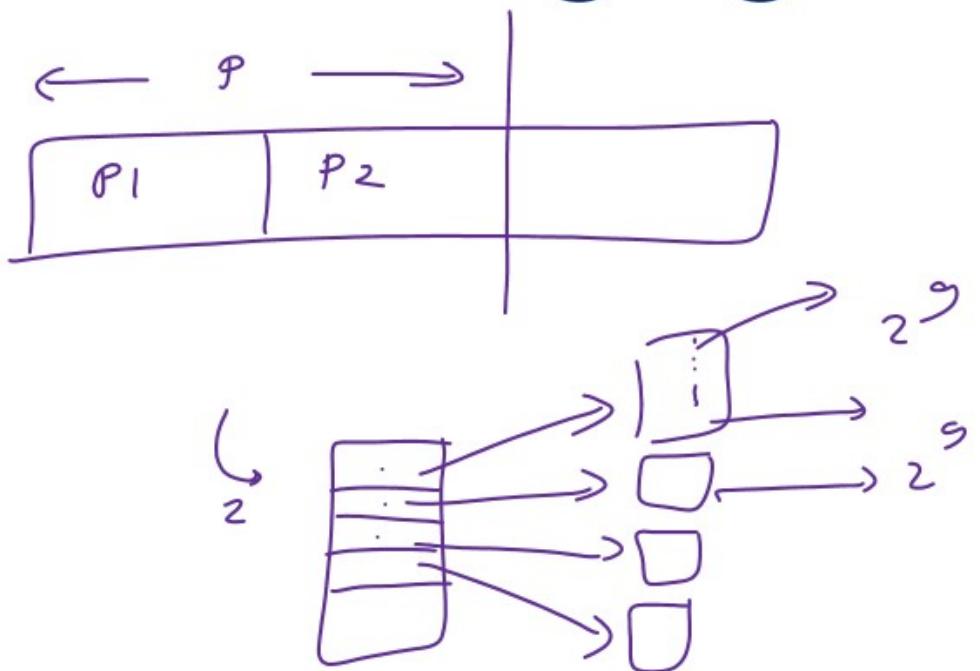
$$\sqrt{2 * \text{Process size} * \text{P.T. entry size}}$$



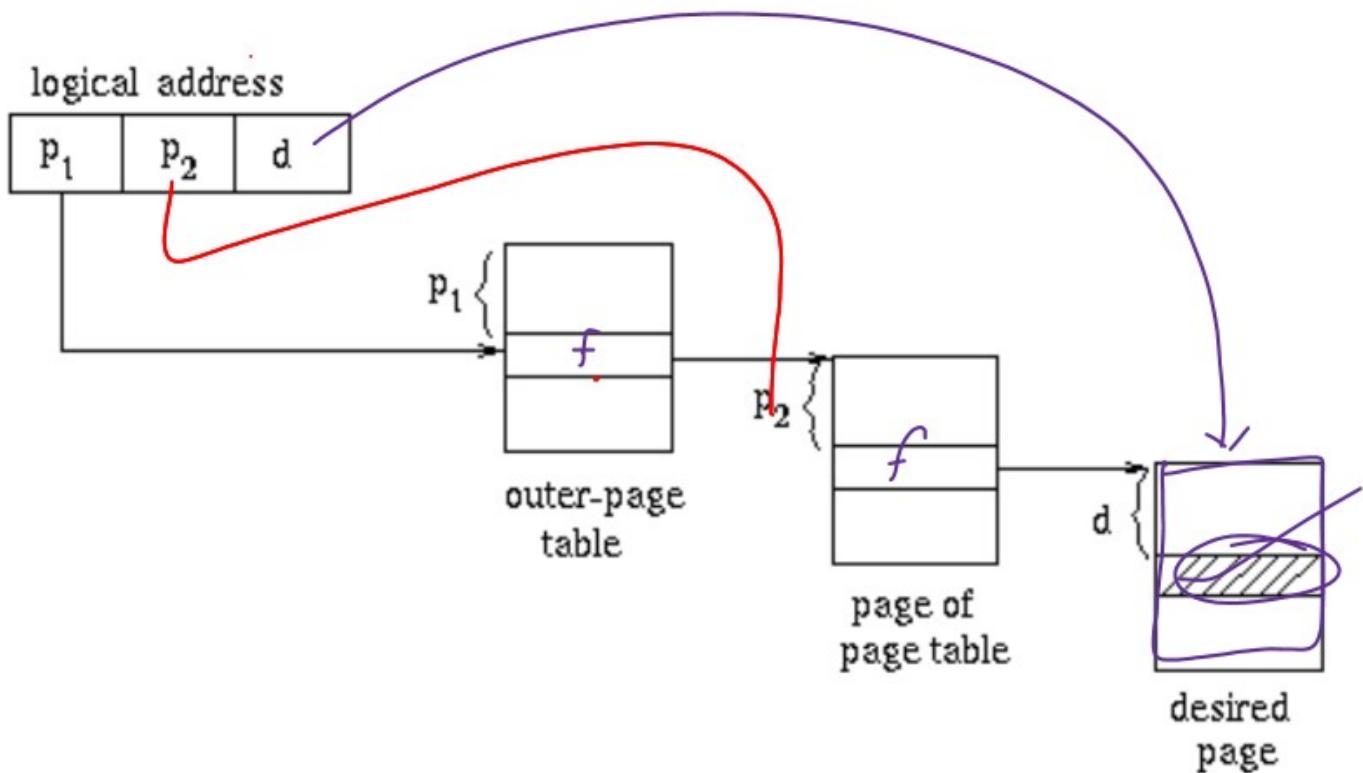
Page Table in Memory



Multilevel Paging



Multilevel Paging



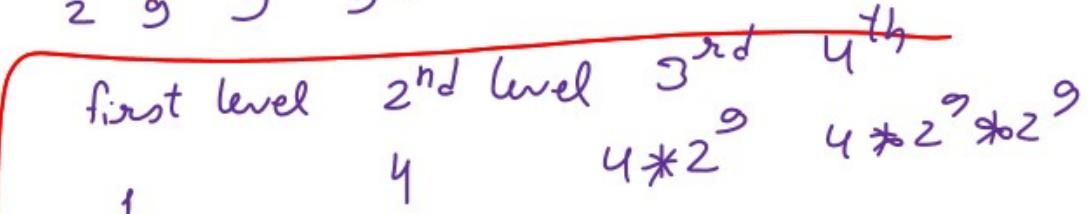
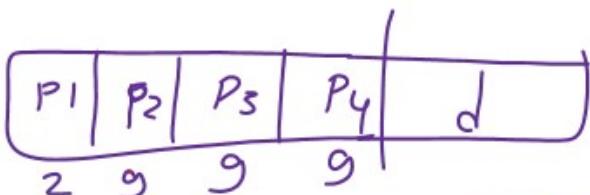
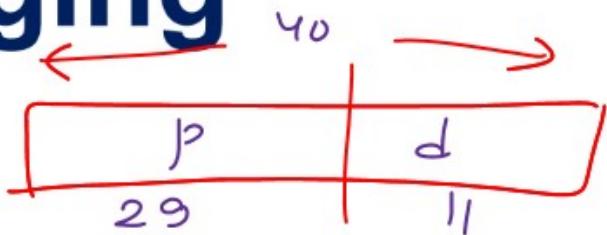
Multilevel Paging

Page size = 2 kB

P.T. ent. = 4 B

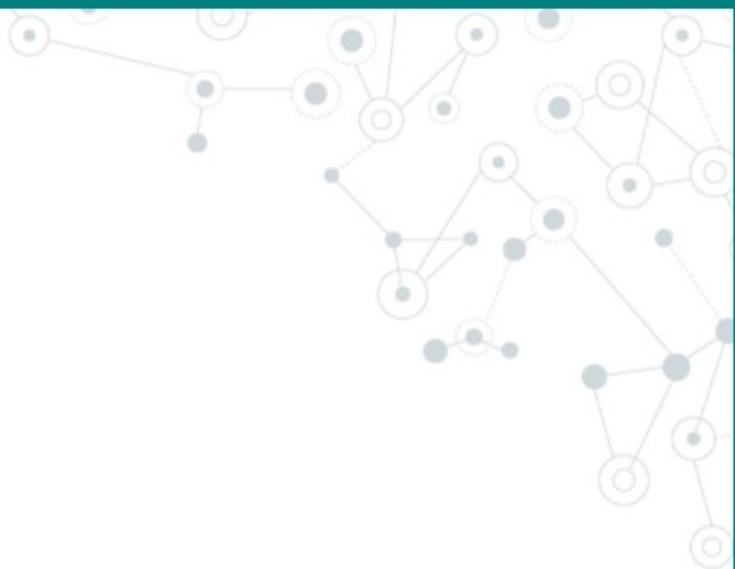
L.A. = 40-bits

$$\text{no. of entries in 1 page} = \frac{2 \text{ kB}}{4 \text{ B}} \\ = 2^9$$

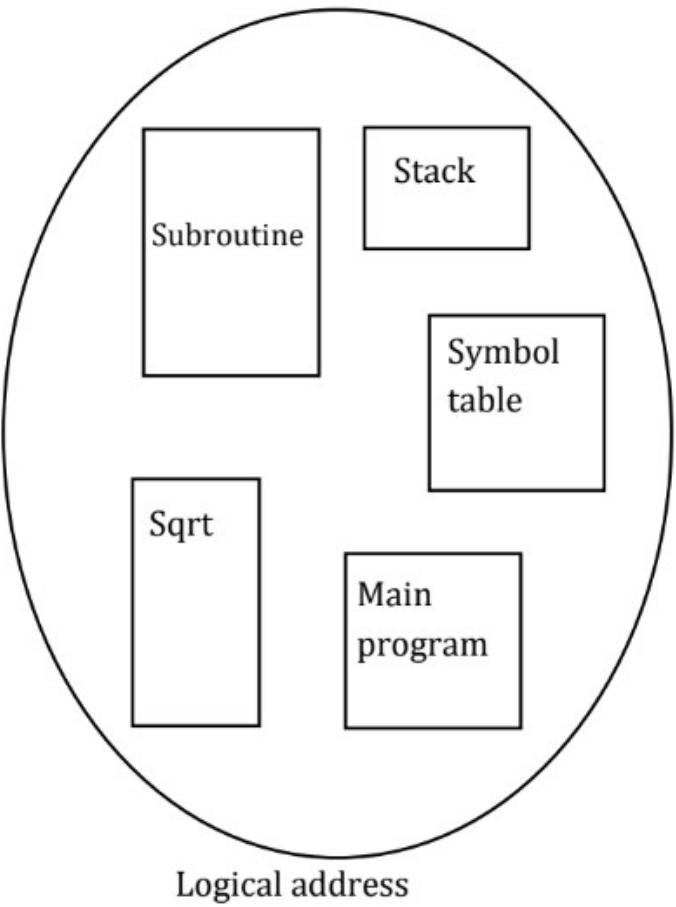


Segmentation

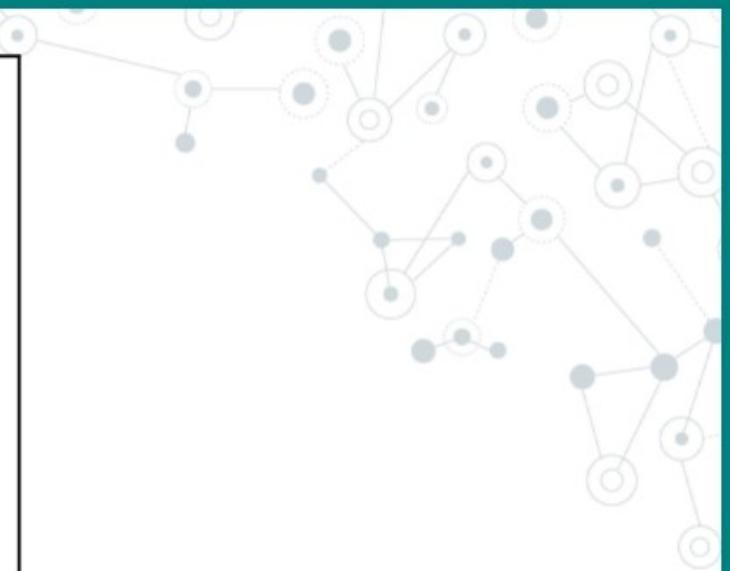
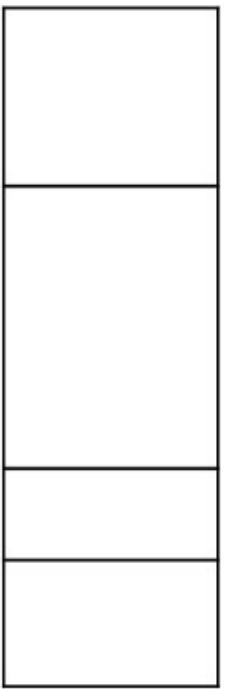
- ➊ Divide Process in logically related partitions (Segments)
- ➋ Segments are scattered in physical memory



Segmentation



Segmentation



Segmentation

- Size of segment can vary, so along with base, keep limit information also
- Limit defines max number of words within the segment



Segmentation

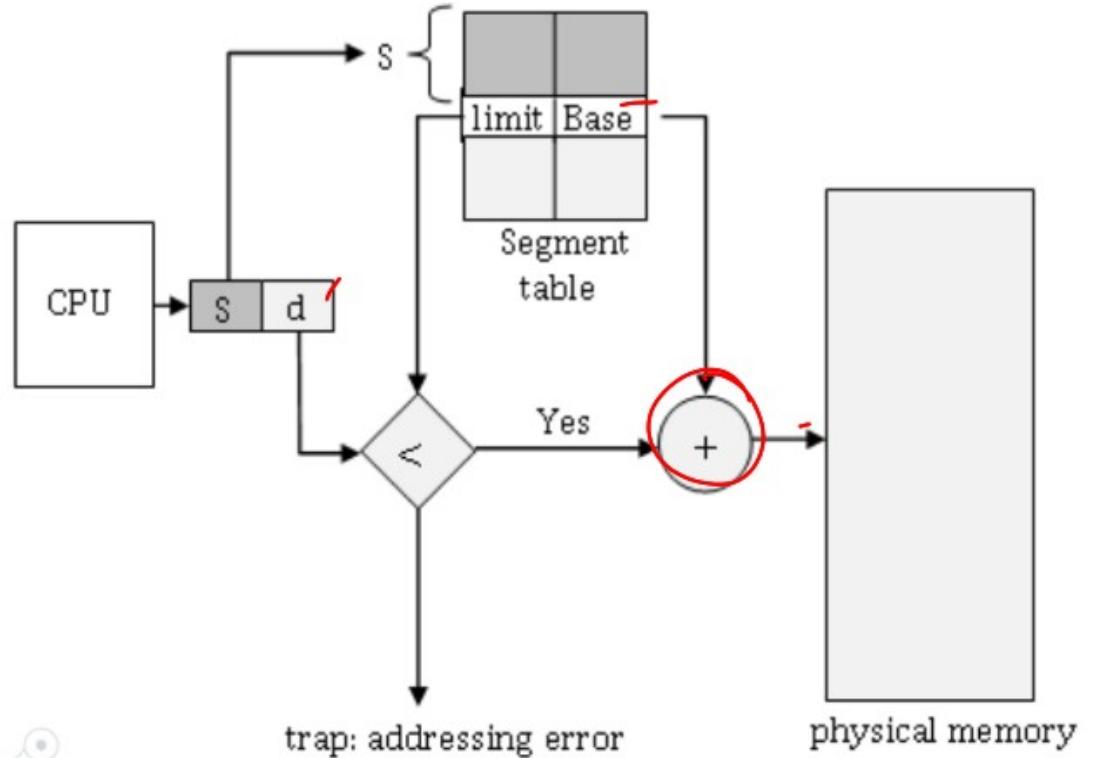
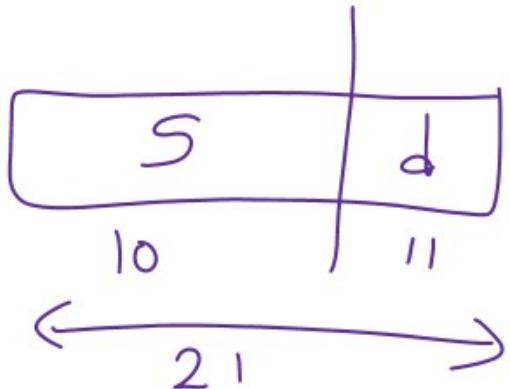


Fig. Segmentation hardware

Question

- Maximum segment size = 2KB
- Number of segments in process = 2^{10}
- Logical address = _____ bits ??

$d = 11$



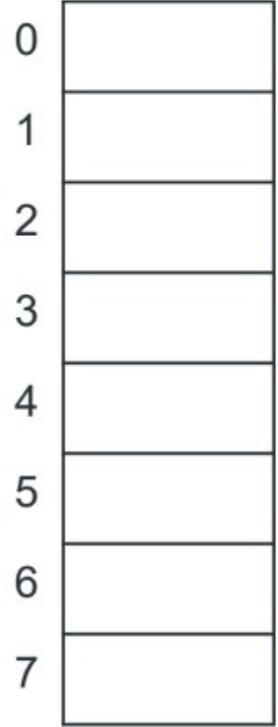
Virtual Memory

- Feature of OS
- Enables to run larger process with smaller available memory



Virtual Memory

Process



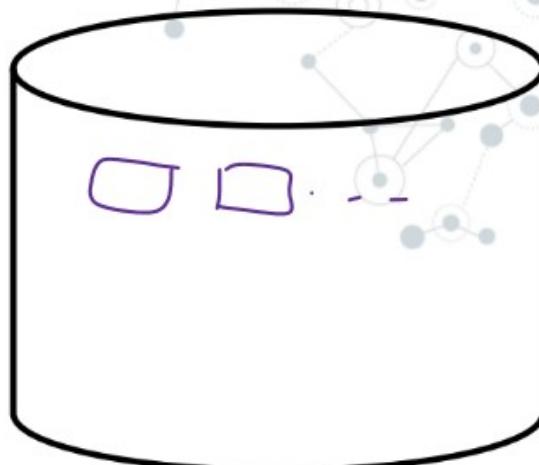
Page Table

	0	1
0	0	0
1	1	1
2	2	1
3	0	0
4	3	1
5	0	1
6	0	0
7	0	0

Physical Memory

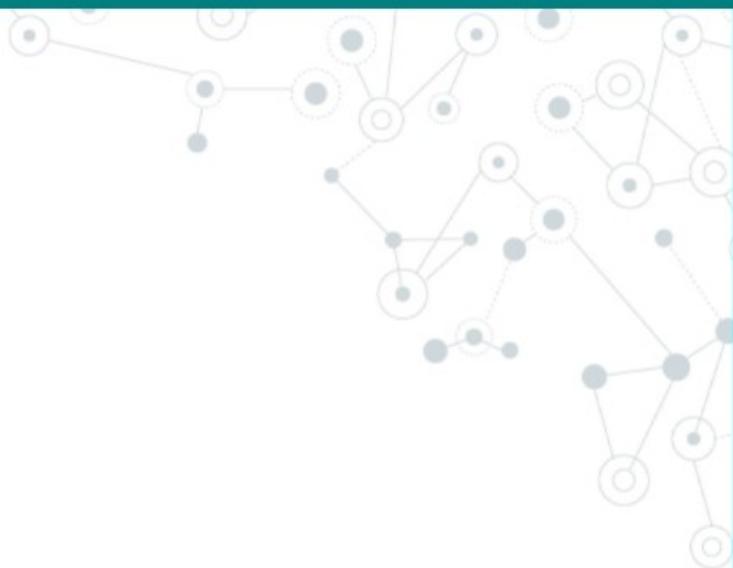
0	Page 6
1	Page 2
2	Page 3
3	Page 5

Secondary Memory



Demand Paging

- ◎ Bring pages in memory when CPU demands



Page Replacement Policies

1. First In First Out (FIFO)
2. Optimal Policy
3. Least Recently Used (LRU)
4. Least Frequently Used (LFU) ✘
5. Most Frequently Used (MFU) ✘
6. Last In First Out (LIFO)



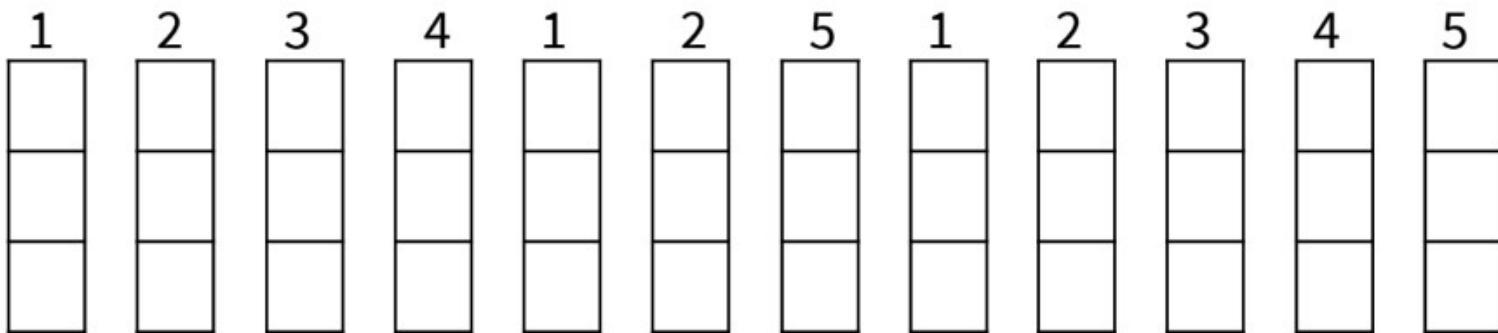
First In First Out (FIFO)

Assume:

- Number of frames = 3 (All empty initially)
- Page reference sequence: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

Page fault rate

$$= \frac{\text{no. of P.f.}}{\text{Total references}}$$



Belady's Anomaly

only in FIFO



First In First Out (FIFO)

Advantages

1. Simple and easy to implement.
2. Low overhead.

Disadvantages:

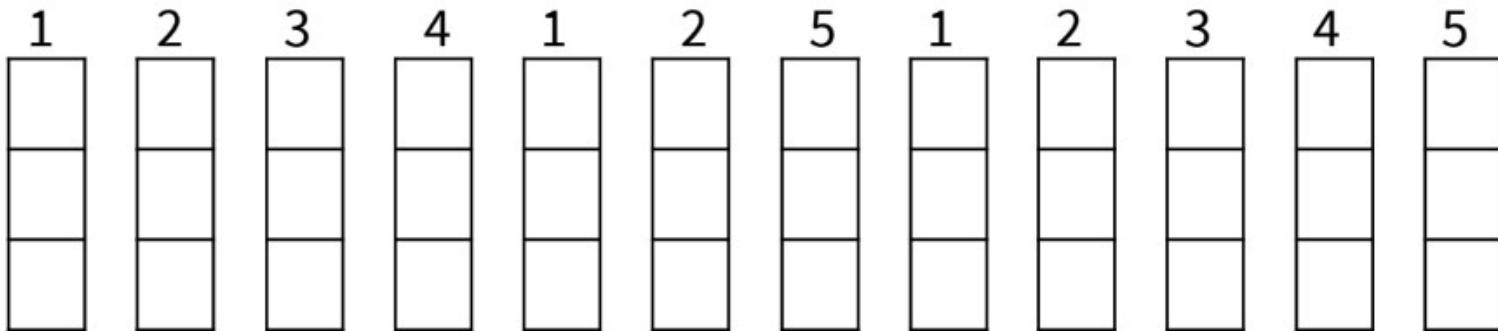
1. Poor performance.
2. Doesn't consider the frequency of use or last used time, simply replaces the oldest page.
3. Suffers from Belady's Anomaly

Optimal Policy

→ replace a page which is going to refer very late in future

Assume:

- Number of frames = 3 (All empty initially)
- Page reference sequence: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



Optimal Policy

Advantages

1. Easy to Implement ✓
2. Simple data structures are used ✓
3. Highly efficient $\Rightarrow \min$ page fault

Disadvantages:

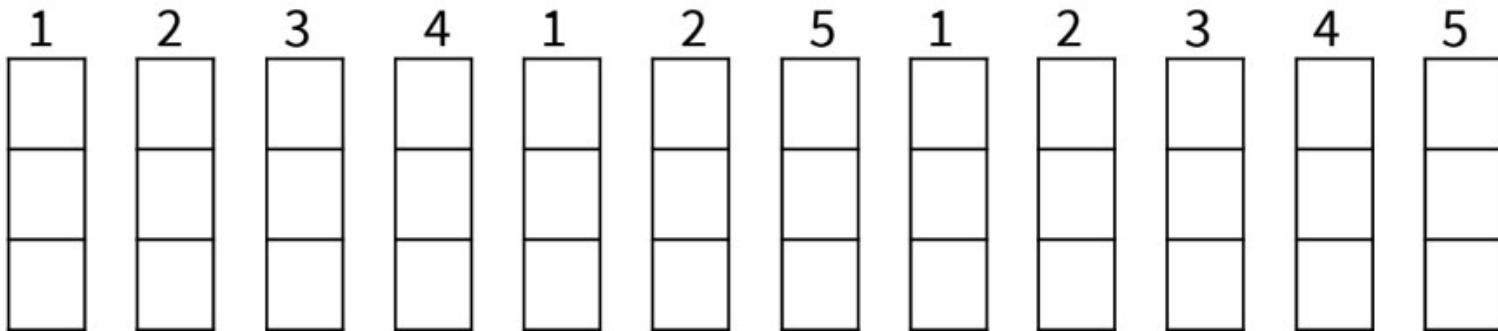
1. Requires future knowledge of the program ✓
2. Time-consuming



Least Recently Used (LRU)

Assume:

- Number of frames = 3 (All empty initially)
- Page reference sequence: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



Least Recently Used (LRU)

Advantages

1. Efficient. ✓
2. Doesn't suffer from Belady's Anomaly ✓

Disadvantages:

1. Complex Implementation ✓
2. Expensive ✓
3. Requires hardware support ✓

Counting Algorithms

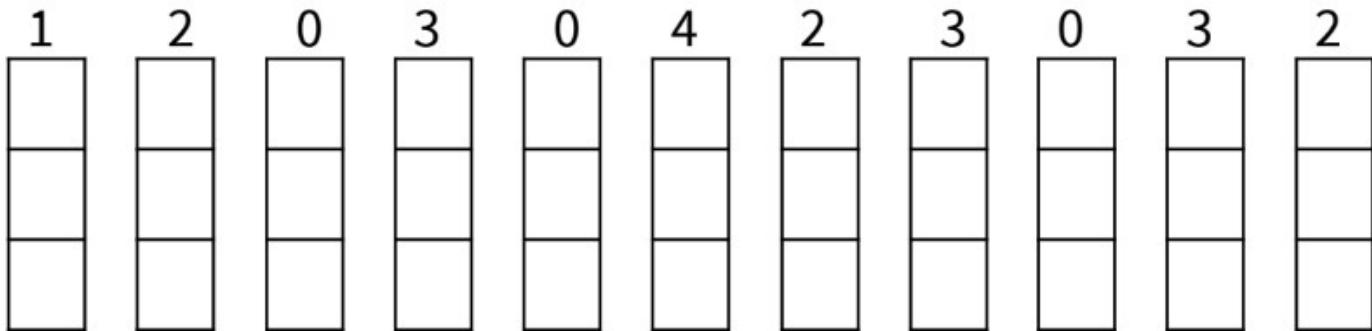
- ◎ Counting algorithms look at the number of occurrences of a particular page and use this as the criterion for replacement.
- ◎ Such counting algorithms includes:
 - LFU (Least Frequently Used)
 - MFU (Most Frequently Used)



Least Frequently Used (LFU)

Assume:

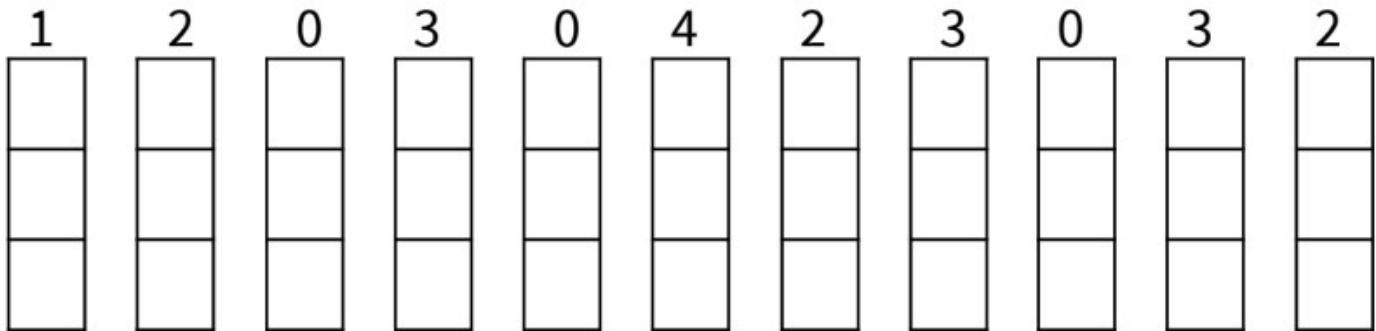
- Number of frames = 3 (All empty initially)
- Page reference sequence: 1 2 0 3 0 4 2 3 0 3 2



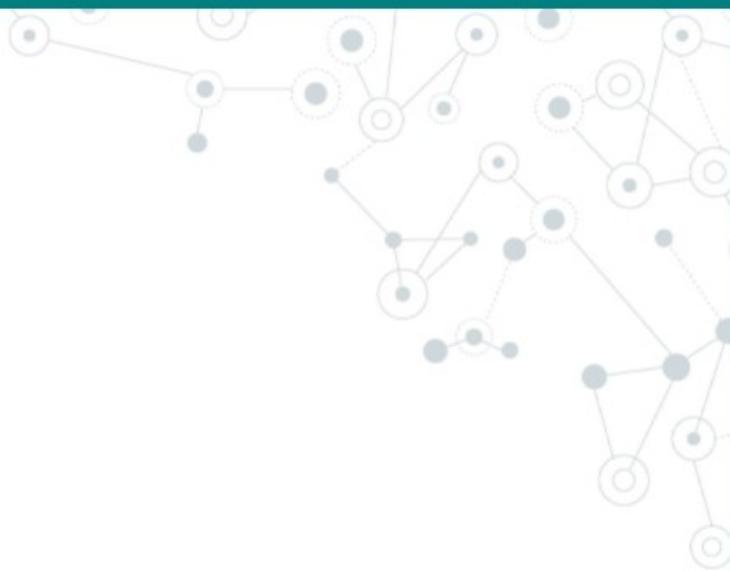
Most Frequently Used (MFU)

Assume:

- Number of frames = 3 (All empty initially)
- Page reference sequence: 1 2 0 3 0 4 2 3 0 3 2

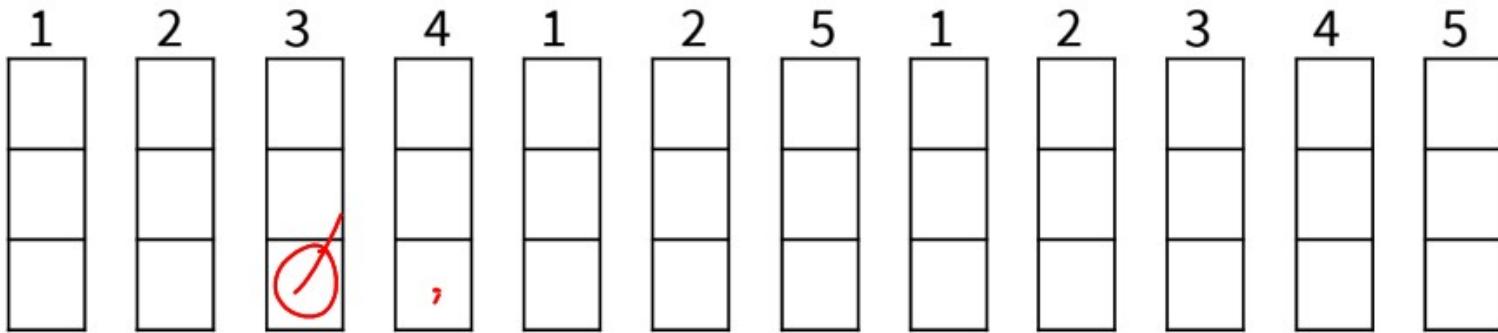


Last In First Out (LIFO)



Assume:

- Number of frames = 3 (All empty initially)
 - Page reference sequence: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



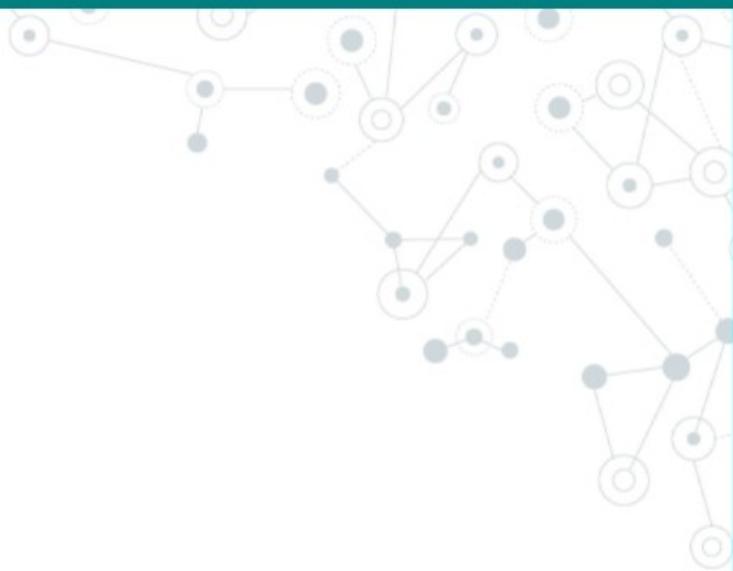
Frame Allocation

How many frames do we allocate per process?

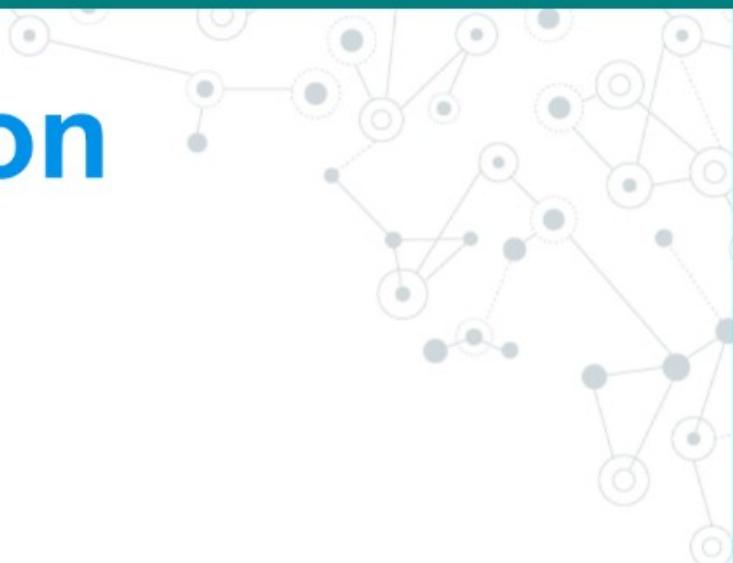
- ◎ If it is a single-user, single-tasking system, it's simple – all the frames belong to the user's process

Frame Allocation

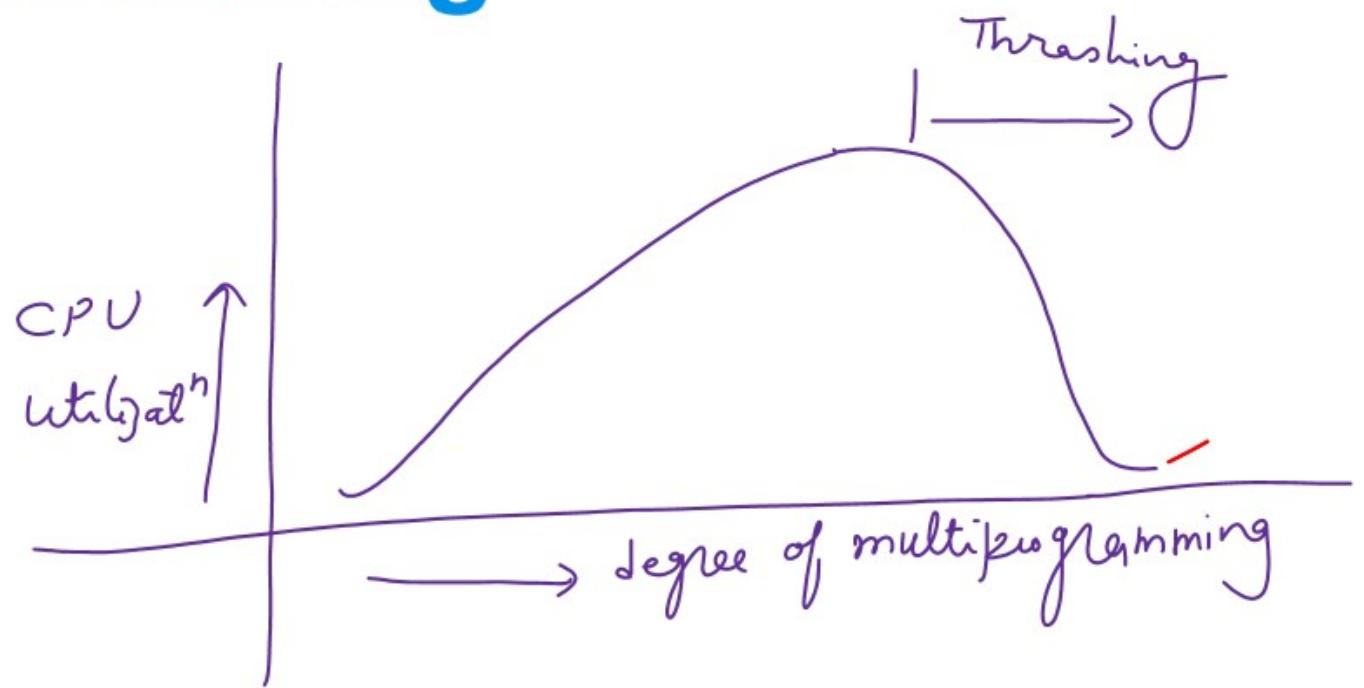
1. Equal Allocation ✓
2. Proportional Allocation ✗



Local Vs Global Allocation

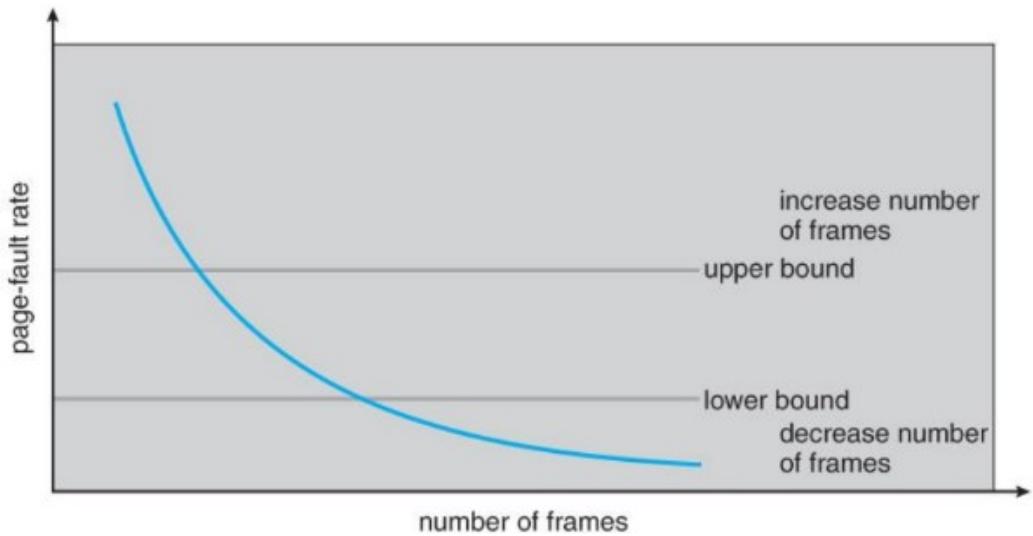


Thrashing



How to Handle Thrashing

1. Working Set Model
2. Page Fault Frequency





Let's Take a Break

See you in 10min



 unacademy

